

# Linux Virtual File System

---

 likegeeks.com/linux-virtual-file-system

The Linux virtual file system or virtual file system generally is a layer that sits on the top of your actual file system which allows the user to access different types of file systems, you can think of virtual file system as an interface between the kernel and the actual file system.

That means you will not find any entries for those Linux virtual filesystems in your `/etc/fstab` file. Yet, you will still find them when you type the `mount` command.

If you are coming from Windows, the virtual file system is the Registry.

## /proc File System

---

The `proc` file system is a virtual file system which is mounted on `/proc` directory.

There is no real file system exists on `/proc`, it's a virtual layer that is used for dealing with the kernel functionalities.

For example, to get the processor specifications, type the following command:

```
$ cat /proc/cpuinfo
```

This is a very powerful and easy way to query Linux kernel.

Notice that if you check the size of the file in `/proc` directory, you will find that all file sizes are 0, because as we said they don't exist on the disk.

When you type **`cat /proc/cpuinfo`** command, a file is dynamically created to show you the CPU info.

The only file that has a size in `/proc` directory is `/proc/kcore` file, which shows the RAM content. Actually, this file isn't occupying any space on the disk.

## Writing to Proc Files

---

As we've seen, we can read the content of `proc` files, but some of them are writable, so we can write to them to change some functionality.

For example, this `/proc/sys/net/ipv4/ip_forward` file controls IP forwarding in case you have multiple network cards.

You can change the value of this file like this:

```
$ echo "1" > /proc/sys/net/ipv4/ip_forward
```

Keep in mind that when you change any file or value under `/proc` directory there is no validation of what you are doing, you may crash your system if you type a wrong setting.

## Persisting /proc Files Changes

---

The previous modification to the **`/proc/sys/net/ipv4/ip_forward`** entry will not survive after rebooting since you are not writing to a file, this is a virtual file system, means change happens to the memory.

If you need to save changes under `/proc`, you have two ways:

- You can write your entries in `/etc/rc.local` file, or in Red Hat based distros like CentOS, create `/etc/rc.d/rc.local` file and make it executable and enable the `systemd` service unit that enables the use of the `rc.local` file and write your entries.
- The `sysctl` command is used to change entries in `/proc/sys/` directory.

```
$ sysctl net.ipv4.ip_forward
```

This will show the value of the entry, to change it, use the `-w` option:

```
$ sysctl -w net.ipv4.ip_forward=1
```

One final step is to write the changes to `/etc/sysctl.conf`:

```
$ echo "net.ipv4.ip_forward = 1" >> /etc/sysctl.conf
```

Make sure that the file `/etc/sysctl.conf` does not contain the entry before you write your changes.

## Common /proc Entries

---

These are some of the commonly used `/proc` entries:

<code>/proc/cpuinfo</code>	information about CPUs in the system.
<code>/proc/meminfo</code>	information about memory usage.
<code>/proc/ioports</code>	list of port regions used for I/O communication with devices.
<code>/proc/mdstat</code>	display the status of RAID disks configuration.
<code>/proc/kcore</code>	displays the system actual memory.
<code>/proc/modules</code>	displays a list of kernel loaded modules.
<code>/proc/cmdline</code>	displays the passed boot parameters.
<code>/proc/swaps</code>	displays the status of swap partitions.
<code>/proc/iomem</code>	the current map of the system memory for each physical device.
<code>/proc/version</code>	displays the kernel version and time of compilation.
<code>/proc/net/dev</code>	displays information about each network device like packets count.
<code>/proc/net/sockstat</code>	displays statistics about network socket utilization.
<code>/proc/sys/net/ipv4/ip_ local_port_range</code>	display the range of ports that Linux uses.
<code>/proc/sys/net/ipv4/ tcp_syncookies</code>	protection against syn flood attacks.

These are some of the common entries in `/proc` directory.

## Listing `/proc` Directory

---

If you list the files in `/proc` directory, you'll notice a lot of directories which have numeric names, these directories contain information about the running processes and the numeric value is the corresponding process ID.

You can check the consumed resources by a specific process from these

directories.

If you take a look at the folder named 1, it belongs to the init process or systemd (like CentOS 7) which is the first process runs When Linux starts.

```
$ ls -l /proc/1
```

---

The /proc/1/exe file is a symbolic link to /lib/systemd/systemd binary or /sbin/init in other systems that use init binary.

The same concept applies to all numeric folders under /proc directory.

## /proc Useful Examples

---

To protect your server from SYN flood attack, you can use iptables to block SYN packets.

A better solution is to use SYN cookies. A special method in the kernel that keeps track of which SYN packets come. If the SYN packets don't move to established state within a reasonable interval, the kernel will drop them.

```
$ sysctl -w net.ipv4.tcp_syncookies=1
```

And to persist the changes.

```
$ echo "net.ipv4.tcp_syncookies = 1" >> /etc/sysctl.conf
```

Another useful example which is the **/proc/sys/fs/file-max**, this value shows the maximum files (including sockets, files, etc,) that can be opened at the same time.

You can increase this number like this:

```
1 $ sysctl -w "fs.file-max=96992"
2 $ echo "fs.file-max = 96992" >>
  /etc/sysctl.conf
3
```

## sysfs Virtual File System

---

sysfs is a Linux virtual file systems which mean it's also in memory.

sysfs file system can be found at **/sys**. The sysfs can be used to get information about your system hardware.

```
$ ls -l /sys
```

From the result of the above command, the file sizes are all zero because as we know this is a Linux virtual file system.

The top level directory of /sys contains the following:

Block	list of block devices detected on the system like sda.
Bus	contains subdirectories for physical buses detected in the kernel.
class	describes class of device like audio, network or printer.
Devices	list all detected devices by the physical bus registered with the kernel.
Module	lists all loaded modules.
Power	the power state of your devices.

## tmpfs Virtual File System

---

tmpfs is a Linux virtual file system that keeps data in the system virtual memory. It is the same like any other Virtual File Systems, any files are temporarily stored in the Kernel's internal caches.

The /tmp file system is used as the storage location for temporary files.

The /tmp file system is backed by an actual disk-based storage and not by a virtual system.

This location is chosen during Linux installation.

The /tmp is created automatically by systemd service when booting the system.

You can setup tmpfs style file system with the size you want, using the mount command.

```
$ mount -t tmpfs -o size=2GB tmpfs /home/myfolder
```

Awesome!!

Working with Linux virtual file system is very easy.

I hope you find the post useful and interesting. Keep coming back.

Thank you.