

How to Prepare for the Confluent Certified Developer for Apache Kafka (CCDAK) exam

M medium.com/@stephane.maarek/how-to-prepare-for-the-confluent-certified-developer-for-apache-kafka-ccdak-exam-ab081994da78

Getting the Apache Kafka certification from Confluent is a great way of making sure to have your skills recognized by your current and future employees. But how do you prepare for that certification?

Certification

I personally find the CCDAK certification of **medium difficulty**, and if you have seriously used Apache Kafka in the past year, many questions shouldn't be surprising. Nonetheless, **it is best to be as prepared as you can**, as the certification cost of **150 USD** is not something you'd rather pay for twice.



I was part of the beta testers :) <https://www.credential.net/6ar0r04w>

I have created a practice exam course with 3 practice tests of 50 questions each... so 150 sample questions. [Check it out here \(discount coupon included!\)](#). Completing those is a great first step to assess where you stand today. **You will find five sample questions in this blog, read on !**

Please also read the CCDK certification FAQ at:
<https://www.confluent.io/certification/>.

Important points are:

- You can currently take the certification straight from your computer, you don't need to go to an exam center
 - The certification exam is 55 multiple choice questions in 90 minutes
 - The certification is valid for two years
 - Read the exam details at: <https://www.confluent.io/wp-content/uploads/CCDAF.pdf>
-

Studying for the Certification

If you find lacking the knowledge, you may already know I have plenty of online content with the [Apache Kafka Series](#) to help you learn Kafka at your own pace and with over 35 hours of videos. **Countless number of my students did successfully pass the CCDAK certification exam after studying with my online courses.**

Let's deep dive into what you need to know for each exam topic:

Apache Kafka Fundamentals

You need to know everything about brokers, topics, partitions, offsets, producers and message keys, consumers and consumer groups, delivery semantics, Zookeeper, Java programming, and more. My [Apache Kafka for Beginners course](#) will help you get there

Kafka Connect

You'll need to know the basics of Kafka Connect includes Workers, Source Connectors and Sink Connectors, as well as converters. [Check out my Kafka Connect Course](#)

Kafka Streams

You will need to know how Kafka Streams works, reading high level DSL topologies and understanding the difference between stateless and stateful operations. [Check out my Kafka Streams course](#)

Confluent Schema Registry / REST Proxy / KSQL

You will need to know the basics of the Confluent components (it is, after all, a Confluent certification), especially the Confluent Schema Registry and how Avro works. [Check out my Confluent Schema Registry course](#). For KSQL, reading a few recipes ahead of time may help: <https://www.confluent.io/stream-processing-cookbook/>. You'll quickly get the idea, it looks like SQL. Reading the FAQ may help too: <https://docs.confluent.io/current/ksql/docs/faq.html>

Kafka Setup

Surprising as it seems, you will get questions about Kafka heap sizing, and setup questions, so you will need to acquire knowledge there too. [Check out my Kafka Setup course](#)

Kafka Security and Kafka Monitoring?

For now, these aren't covered at the exam in depth. My blog on [Introduction to Apache Kafka Security](#) will do.

Sample questions

The following questions are an extract of my [Practice Exam course](#) (responses at the bottom of the blog):

Q1: You want to perform table lookups against a KTable every time a new record is received from a KStream. What is the output of KStream-KTable join?

- A: KStream
- B: KTable
- C: GlobalKTable
- D: You choose between KStream or KTable

Q2: In Kafka, every broker... (select three)

- A: contains all the topics and all the partitions
- B: contains only a subset of the topics and the partitions
- C: knows all the metadata for all topics and partitions
- D: knows the metadata for the topics and partitions it has on its disk
- E: is a bootstrap broker
- F: is a controller

Q3: If a topic has a replication factor of 3...

- A: Each partition will live on 3 different brokers
- B: Each partition will live on 2 different brokers
- C: Each partition will live on 4 different brokers
- D: 3 replicas of the same data will live on 1 broker

Q4: Your manager would like to have topic availability over consistency. Which setting do you need to change in order to enable that?

- A: min.insync.replicas
- B: compression.type
- C: unclean.leader.election.enable

Q5: Using the Confluent Schema Registry, where are Avro schema stored?

- A: In the Schema Registry embedded SQL database
- B: In the `_schemas` topic
- C: In the message bytes themselves
- D: In the Zookeeper node `/schemas`

Closing Comments

Don't rush your certification, especially if you're just getting started with Apache Kafka, as I find the exam really great at testing your actual experience using Apache Kafka. Take your time to learn and practice, and answering the questions will seem natural.

If I have helped you get the certification, leave a comment here or [a review on my Udemy course!](#)

Happy learning, and I wish you good luck for the CCDAK certification!

Quiz answers: 1A, 2BCE, 3A, 4C, 5B.